

## Problem 1. Implementing a Neural Network Classifier

In this homework, you will implement a neural network and the backpropagation algorithm and stochastic gradient descent with mini-batches. The goal is to implement this yourself with only the basic linear algebra package *numpy*, rather than a modern deep learning library. We will use MNIST, a common classification data set consisting of 28x28 pixel grayscale images of handwritten digits (0-9), with 60,000 training examples and 10,000 test examples. The task is to probabilistically predict the digit from the image using multi-class classification.




Figure 1: MNIST example images

Requirements:

- (a) Your model should use a *softmax* output layer with categorical cross-entropy loss.

$$\text{Loss}_{ce} = \sum_i^C y_i \log(\hat{y}_i) \quad (1)$$

- (b) Your model should have three hidden layers. The layers should be fully connected. The activation/transfer function of the hidden layers should have ReLU, SELU, and hyperbolic tangent functions, respectively.

$$\text{ReLU}(x) = \max(0, x) \quad (2)$$

$$\text{SELU}(x) = \begin{cases} 1.75809 * (e^x - 1) & \text{for } x \leq 0 \\ 1.0507 * x & \text{for } x > 0 \end{cases} \quad (3)$$

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}} \quad (4)$$

- (c) Use mini-batch gradient descent (e.g. mini-batches of size 100).
- (d) Implement momentum, RMSProp, or ADAM optimization algorithm.
- (e) Train the network to achieve at least 94% accuracy on the test set. Tune the hyper-parameters (e.g. weight initialization, learning rate) to get good performance on the test data.

- (f) Plot both the training and test *loss* (the NLL) throughout training (at every epoch).
- (g) Write a paragraph describing the network, training algorithm, and any hyperparameter choices. Make sure to report the *accuracy* (also called the *01-loss*. To get a class prediction from a softmax layer, we simply take the *argmax* of the probabilities.
- (h) Plot another figure that demonstrates *overfitting* to the training data, so that the test loss increases while the training loss continues to decline. (Hint: this should be easier with a network of a single hidden layer, so you can change the architecture.)
- (i) Submit your homework as a Jupyter notebook to the github classroom.

Tips:

- (a) For a minimalist implementation of a neural network, see:  
<https://iamtrask.github.io/2015/07/12/basic-python-network/>
- (b) To quickly download the MNIST data set, see:  
<https://github.com/hsjeong5/MNIST-for-Numpy>
- (c) A single pass through the training data is called an *epoch*. You may have to do 50-100 epochs to see overfitting. If your computer is slow, go ahead and use a fraction of the data set, e.g. 6,000 examples.
- (d) We typically compute the validation loss at the end of every training epoch to see if we are overfitting. For the training data set, we usually keep a running average of the loss while performing the mini-batch training updates to save on computation.